

Getting Started with Pandas

March 31, 2015

1 Pandas in a nutshell

Pandas is the Python library designed specifically for data analysis. The author of *Python for Data Analysis*, Wes McKinney, began developing **Pandas** in 2008

while at [AQR Capital Management](#) out of need for a performant, flexible tool to perform quantitative analysis on financial data. Before leaving AQR he was able to convince management to allow him to open source the library.

Another AQR employee, Chang She, joined the effort in 2012 as the second major contributor to the library. Right around that time, the library became popular in the Python community, and many more contributors joined the project making it one of the most vital and active data analysis libraries for Python. ([Wikipedia](#))

Pandas can be thought of the Python equivalent of Microsoft Excel. It abstracts the notion of the spreadsheet, allowing the user to use powerful and robust analytical tools generally to automate repeated processes.

The twin centerpieces of the **Pandas** library are the **Series** and the **DataFrame**. The **Series** class is, at its core, a one-dimensional NumPy array, surrounded by additional information, such as its index. The **DataFrame** is conceptually an *array* of **Series** classes, each sharing the same index.

```
In [1]: from pandas import Series, DataFrame
import pandas as pd
```

We will be using Wes McKinney's [GitHub notebook](#) as a skeleton. He imports the following libraries for later use:

```
In [2]: from __future__ import division
from numpy.random import randn
import numpy as np
import os
import matplotlib.pyplot as plt
np.random.seed(12345)
plt.rc('figure', figsize=(10, 6))
from pandas import Series, DataFrame
import pandas as pd
np.set_printoptions(precision=4)
```

2 Introduction to pandas data structures

2.1 Series

Consider the following input:

```
In [3]: obj = Series([4, 7, -5, 3])
        obj
```

```
Out[3]: 0    4
        1    7
        2   -5
        3    3
        dtype: int64
```

We have set the variable `obj` to reference a new **Pandas Series**, which we initialized by giving a **Python** list as input. Notice that **Pandas** automatically interprets the input data as type `int64`, which indicates that it is fairly smart! Also, notice that upon printing `obj` we see *two* columns. The first column is the *index* of the **Series** class, which is presently the natural index, `range(4)`. The second column is the input data that we gave initially, which **Pandas** refers to as the *values* of `obj`. You can access these columns individually by calling `obj.index` and `obj.values`, respectively. For example:

```
In [4]: print obj.index, "\n\n"
        print obj.values

Int64Index([0, 1, 2, 3], dtype='int64')
```

```
[ 4  7 -5  3]
```

2.1.1 Indexing

As previously mentioned, the natural index simply starts at 0 and increments integers to the size of the list of the input values. Alternatively, we can specify the index explicitly when we initialize the **Series**, as in the following:

```
In [5]: obj2 = Series([4, 7, -5, 3], index=['d', 'b', 'a', 'c'])
        obj2
```

```
Out[5]: d    4
        b    7
        a   -5
        c    3
        dtype: int64
```

What happens when you examine the index now?

```
In [6]: obj2.index

Out[6]: Index([u'd', u'b', u'a', u'c'], dtype='object')
```

Don't be alarmed by the `u` prefix to each of the index values. In **Python** (as well as in other languages, this simply indicates a **Unicode string**. Ostensibly, there is no difference between normal strings and Unicode strings.

You can access a particular member of the **Series** data by specifying its index. For example,

```
In [7]: obj2['a']

Out[7]: -5
```

This allows you to change the value of specific entries in your **Series** data. Additionally, you can call *subSeries* by specifying a sublist of the index.

```
In [8]: obj2['d'] = 6
        obj2[['c', 'a', 'd']]
```

```
Out[8]: c    3
        a   -5
        d    6
        dtype: int64
```

A powerful tool in **Pandas** is the ability to concisely access data meeting **Boolean** qualifications. In the case below, `obj2 > 0` is given as the “index,” and the output is the **subSeries** of `obj2` for which all entries are positive.

```
In [9]: obj2[obj2 > 0]
```

```
Out[9]: d    6
        b    7
        c    3
        dtype: int64
```

2.1.2 Aside

What is `obj2 > 0` actually?

```
In [10]: obj2 > 0
```

```
Out[10]: d    True
         b    True
         a   False
         c    True
         dtype: bool
```

This is actually a neat property of **Pandas** which is similar to **NumPy**. In **NumPy**, suppose you are given an array:

```
In [11]: arr = np.random.rand(5) * 2.0 - 1.0
```

The actual array itself is given by

```
In [12]: arr
```

```
Out[12]: array([ 0.8592, -0.3672, -0.6322, -0.5909,  0.1355])
```

The **Boolean** array specifying which elements of `arr` are positive is given by

```
In [13]: boolArr = arr > 0.0
        boolArr
```

```
Out[13]: array([ True, False, False, False,  True], dtype=bool)
```

Similarly, you can generate a new **Series** of **Boolean** values by subjecting the original **Series** to a **Boolean** statement, as we did above.

2.1.3 Broadcasting

Like NumPy, we can *broadcast* arithmetic operations onto **Series** data. For example,

```
In [14]: obj2 * 2
```

```
Out[14]: d    12
         b    14
         a   -10
         c     6
         dtype: int64
```

returns a **Series** whose values are doubled, and

```
In [15]: np.exp(obj2)
```

```
Out[15]: d    403.428793
         b   1096.633158
         a     0.006738
         c    20.085537
         dtype: float64
```

returns a **Series** whose values have all been subject to the transformation $x \mapsto e^x$. Notice additionally that the **dtype** of `obj2` has automatically been changed from **int64** to **float64**. Again, **Pandas** is being smart!

2.1.4 Querying a Series

In Python, there is a binary operator called **in**, which takes two “arguments.” The left-hand argument is can be any type of data (or object, we won’t get into this), while the right-hand argument is some type of iterable object. Then **in** returns **True** if the left-hand argument is an *element* of the right-hand argument. Mathematically, this is equivalent to set membership. For example,

```
In [16]: odds = [i for i in range(20) if i%2 == 1]
         print 3 in odds, "|", 2 in odds
```

```
True | False
```

is equivalent to noting that if

$$\text{Odds} = \{n : 0 \leq n < 20 \text{ and } n \text{ is odd}\}$$

we have that

$$3 \in \text{Odds}$$

while

$$2 \notin \text{Odds}.$$

(In fact, we have already seen **in** in action with Python’s **for** loop, which has the form

```
for element in iterative_object:
```

indicating that the code should loop through every element **element** that is a member of **iterative_object**.)

You can use **in** with **Pandas Series** to test that an element is a member of the index of the **Series**. For example,

```
In [17]: 'b' in obj2
```

```
Out[17]: True
```

```
In [18]: 'e' in obj2
```

```
Out[18]: False
```

2.1.5 Aside

We have talked about the native Python list data type. There is another important native data type in Python, called a dict, which you can learn about more [here](#). Python dict types are similar to association lists in Scheme, in that they require a lookup key in order to access elements.

Crucially, Pandas can create a Series from a dict by interpreting the key for each item as its corresponding index value, which is actually quite natural. In this sense, I find that it is useful to think of the relationship between NumPy and Pandas as akin to the relationship between a list and a dict.

```
In [19]: sdata = {'Ohio': 35000, 'Texas': 71000, 'Oregon': 16000, 'Utah': 5000}
         obj3 = Series(sdata)
         obj3
```

```
Out[19]: Ohio      35000
         Oregon    16000
         Texas     71000
         Utah      5000
         dtype: int64
```

What happens when you use an existing dataset with a new index, in which there is a new, unfilled index?

```
In [20]: states = ['California', 'Ohio', 'Oregon', 'Texas']
         obj4 = Series(sdata, index=states)
         obj4
```

```
Out[20]: California    NaN
         Ohio          35000
         Oregon        16000
         Texas         71000
         dtype: float64
```

In this case, California is a previously-unused index, which has no corresponding value. Thus, Pandas initializes the new Series with the value corresponding to California set to NaN (Python-speak for null).

The `isnull` method returns a Series of Boolean values whenever the original Series has a null (NaN) value.

```
In [21]: pd.isnull(obj4)
```

```
Out[21]: California    True
         Ohio          False
         Oregon        False
         Texas         False
         dtype: bool
```

The `notnull` method does the exact opposite!

```
In [22]: pd.notnull(obj4)
```

```
Out[22]: California    False
         Ohio          True
         Oregon        True
         Texas         True
         dtype: bool
```

The methods `isnull` and `notnull` are “static” in the sense that they can be called straight from the `pd` module or for a specific Series object.

```
In [23]: obj4.isnull()
```

```
Out[23]: California    True
         Ohio          False
         Oregon         False
         Texas          False
         dtype: bool
```

Recall the two `Series`, `obj3` and `obj4`:

```
In [24]: print "\tobj3:\n",obj3, "\n\n\tobj4:\n", obj4
```

```
obj3:
Ohio      35000
Oregon     16000
Texas      71000
Utah       5000
dtype: int64
```

```
obj4:
California  NaN
Ohio        35000
Oregon      16000
Texas       71000
dtype: float64
```

Arithmetic operations between distinct `Series` objects work conservatively. For data types, `int64 + float64 = float64` to preserve the decimal information. The summed index is the union of the two indices. Consider the following example:

```
In [25]: obj3 + obj4
```

```
Out[25]: California    NaN
         Ohio          70000
         Oregon         32000
         Texas         142000
         Utah           NaN
         dtype: float64
```

No entry for `California` exists in `obj3`, while no entry for `Utah` exists in `obj4`. Pandas interprets `NaN + x = NaN` for all `x`, so the resultant `Series` sets `NaN` for both `California` and `Utah`.

We can set some metadata for a `Series`, such as the name of the values column and the name of the index column.

```
In [26]: obj4.name = 'population'
         obj4.index.name = 'state'
         obj4
```

```
Out[26]: state
         California    NaN
         Ohio          35000
         Oregon        16000
         Texas         71000
         Name: population, dtype: float64
```

You can also completely change the index at any time. This is something we will get into more detail later.

```
In [27]: obj.index = ['Bob', 'Steve', 'Jeff', 'Ryan']
obj
```

```
Out[27]: Bob      4
         Steve    7
         Jeff    -5
         Ryan     3
         dtype: int64
```

2.2 DataFrame

Like we said before, you can think of a `DataFrame` as an *array* of `Series` objects. Specifically, a `DataFrame` is a two-dimensional array of `Series` objects, all indexed by the same index series. You can also think of a `DataFrame` as a single Microsoft Excel spreadsheet.

One way to initialize a `DataFrame` is by giving a `dict` where each key indicates a Python list.

```
In [28]: data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada'],
                 'year': [2000, 2001, 2002, 2001, 2002],
                 'pop': [1.5, 1.7, 3.6, 2.4, 2.9]}
frame = DataFrame(data)

frame
```

```
Out[28]:   pop  state  year
0  1.5   Ohio  2000
1  1.7   Ohio  2001
2  3.6   Ohio  2002
3  2.4  Nevada  2001
4  2.9  Nevada  2002
```

You can reorder the columns in a new `DataFrame` using the following argument:

```
In [29]: DataFrame(data, columns=['year', 'state', 'pop'])
```

```
Out[29]:   year  state  pop
0  2000   Ohio  1.5
1  2001   Ohio  1.7
2  2002   Ohio  3.6
3  2001  Nevada  2.4
4  2002  Nevada  2.9
```

Similarly, the `index` optional argument in `DataFrame` allows you to specify the index list. Additionally, adding a `debt` column with no corresponding data in `data` will initialize a column filled with `NaN` entries.

```
In [30]: frame2 = DataFrame(data, columns=['year', 'state', 'pop', 'debt'],
                           index=['one', 'two', 'three', 'four', 'five'])
frame2
```

```
Out[30]:   year  state  pop  debt
one  2000   Ohio  1.5  NaN
two  2001   Ohio  1.7  NaN
three 2002   Ohio  3.6  NaN
four  2001  Nevada  2.4  NaN
five  2002  Nevada  2.9  NaN
```

You can access the columns of a `DataFrame` as follows:

```
In [31]: frame2.columns
```

```
Out[31]: Index([u'year', u'state', u'pop', u'debt'], dtype='object')
```

You can slice a particular column by specifying its column name. Notice how this returns a **Series**.

```
In [32]: frame2['state']
```

```
Out[32]: one      Ohio
         two      Ohio
         three    Ohio
         four    Nevada
         five    Nevada
         Name: state, dtype: object
```

Alternatively, you can slice a column using the following syntax:

```
In [33]: frame2.year
```

```
Out[33]: one      2000
         two      2001
         three    2002
         four      2001
         five      2002
         Name: year, dtype: int64
```

To slice a row, you can specify an index, which will return a **Series** representing the row at the index.

```
In [34]: frame2.ix['three']
```

```
Out[34]: year      2002
         state    Ohio
         pop       3.6
         debt      NaN
         Name: three, dtype: object
```

Broadcasting works in the natural way that you might expect:

```
In [35]: frame2['debt'] = 16.5
         frame2
```

```
Out[35]:
```

	year	state	pop	debt
one	2000	Ohio	1.5	16.5
two	2001	Ohio	1.7	16.5
three	2002	Ohio	3.6	16.5
four	2001	Nevada	2.4	16.5
five	2002	Nevada	2.9	16.5

You can also give a particular column a list or **ndarray**, which will then be distributed across the column.

```
In [36]: frame2['debt'] = np.arange(5.)
         frame2
```

```
Out[36]:
```

	year	state	pop	debt
one	2000	Ohio	1.5	0
two	2001	Ohio	1.7	1
three	2002	Ohio	3.6	2
four	2001	Nevada	2.4	3
five	2002	Nevada	2.9	4

Finally, you can give a column of a `DataFrame` a `Series`. If you specify a `Series` with an index differing from the main `DataFrame`, then the entries of the `DataFrame` will be set to `NaN`.

```
In [37]: val = Series([-1.2, -1.5, -1.7], index=['two', 'four', 'five'])
         frame2['debt'] = val
         frame2
```

```
Out[37]:
```

	year	state	pop	debt
one	2000	Ohio	1.5	NaN
two	2001	Ohio	1.7	-1.2
three	2002	Ohio	3.6	NaN
four	2001	Nevada	2.4	-1.5
five	2002	Nevada	2.9	-1.7

The point of `Pandas` is there are *numerous* ways to achieve the same effect, depending on whatever is easiest for the task at hand. Here is another way to add a column:

```
In [38]: frame2['eastern'] = frame2.state == 'Ohio'
         frame2
```

```
Out[38]:
```

	year	state	pop	debt	eastern
one	2000	Ohio	1.5	NaN	True
two	2001	Ohio	1.7	-1.2	True
three	2002	Ohio	3.6	NaN	True
four	2001	Nevada	2.4	-1.5	False
five	2002	Nevada	2.9	-1.7	False

We can also use Python's `del` function to remove a column:

```
In [39]: del frame2['eastern']
         frame2.columns
```

```
Out[39]: Index([u'year', u'state', u'pop', u'debt'], dtype='object')
```

One final way to initialize `DataFrame` objects is with nested `dict` objects.

```
In [40]: pop = {'Nevada': {2001: 2.4, 2002: 2.9},
               'Ohio': {2000: 1.5, 2001: 1.7, 2002: 3.6}}

         frame3 = DataFrame(pop)
         frame3
```

```
Out[40]:
```

	Nevada	Ohio
2000	NaN	1.5
2001	2.4	1.7
2002	2.9	3.6

You can transpose a `DataFrame` if it makes more sense to work with the rows and columns flipped.

```
In [41]: frame3.T
```

```
Out[41]:
```

	2000	2001	2002
Nevada	NaN	2.4	2.9
Ohio	1.5	1.7	3.6

You can do this transpose operation from the outset by manually specifying the index.

```
In [42]: DataFrame(pop, index=[2001, 2002, 2003])
```

```
Out[42]:
```

	Nevada	Ohio
2001	2.4	1.7
2002	2.9	3.6
2003	NaN	NaN

DataFrame objects can also be initialized from dicts of Series objects.

```
In [43]: pdata = {'Ohio': frame3['Ohio'][:-1],
                  'Nevada': frame3['Nevada'][:2]}
          DataFrame(pdata)
```

```
Out[43]:
```

	Nevada	Ohio
2000	NaN	1.5
2001	2.4	1.7

```
In [44]: frame3.index.name = 'year'; frame3.columns.name = 'state'
          frame3
```

```
Out[44]:
```

state	Nevada	Ohio
year		
2000	NaN	1.5
2001	2.4	1.7
2002	2.9	3.6

If you need to access the underlying ndarray from any DataFrame, use the DataFrame.values field.

```
In [45]: frame3.values
```

```
Out[45]: array([[ nan,  1.5],
                [ 2.4,  1.7],
                [ 2.9,  3.6]])
```

```
In [46]: frame2.values
```

```
Out[46]: array([[2000, 'Ohio', 1.5, nan],
                [2001, 'Ohio', 1.7, -1.2],
                [2002, 'Ohio', 3.6, nan],
                [2001, 'Nevada', 2.4, -1.5],
                [2002, 'Nevada', 2.9, -1.7]], dtype=object)
```

2.3 Index objects

The Index is the “metadata” object for Series and DataFrame objects. We’ve seen ways of initializing Index objects before, so we will go over some features of these objects.

```
In [47]: obj = Series(range(3), index=['a', 'b', 'c'])
          index = obj.index
          index
```

```
Out[47]: Index([u'a', u'b', u'c'], dtype='object')
```

Index objects can be sliced like arrays.

```
In [48]: index[1:]
```

```
Out[48]: Index([u'b', u'c'], dtype='object')
```

Importantly, Index objects are not mutable, so you can’t change their values in the natural way:

```
In [49]: index[1] = 'd'
```

```
-----

TypeError                                Traceback (most recent call last)

<ipython-input-49-676fdeb26a68> in <module>()
----> 1 index[1] = 'd'

/home/alethiometryst/anaconda/lib/python2.7/site-packages/pandas/core/index.pyc in __setitem__(self, key, value)
894
895     def __setitem__(self, key, value):
--> 896         raise TypeError("Indexes does not support mutable operations")
897
898     def __getitem__(self, key):

TypeError: Indexes does not support mutable operations
```

You can initialize Index objects with NumPy ndarray objects.

```
In [50]: index = pd.Index(np.arange(3))
         obj2 = Series([1.5, -2.5, 0], index=index)
         obj2.index is index
```

```
Out[50]: True
```

```
In [51]: frame3
```

```
Out[51]: state  Nevada  Ohio
         year
         2000      NaN    1.5
         2001      2.4    1.7
         2002      2.9    3.6
```

```
In [52]: print 'Ohio' in frame3.columns, "|", 2003 in frame3.index
```

```
True | False
```

3 Essential functionality

Now that we are familiar with the basic objects in Pandas, we will start working with the mechanics of these objects.

3.1 Reindexing

In the previous section we mentioned that Index objects are immutable. Here we will address this issue.

```
In [53]: obj = Series([4.5, 7.2, -5.3, 3.6], index=['d', 'b', 'a', 'c'])
         obj
```

```
Out[53]: d    4.5
         b    7.2
         a   -5.3
         c    3.6
         dtype: float64
```

The simplest way to change an Index object in an existing Series or DataFrame is with the `reindex` method.

```
In [54]: obj2 = obj.reindex(['a', 'b', 'c', 'd', 'e'])
         obj2
```

```
Out[54]: a   -5.3
         b    7.2
         c    3.6
         d    4.5
         e    NaN
         dtype: float64
```

In the above example, since "e" was not in the original Index, the corresponding Series value is set to NaN. If you want to change the default fill value, `reindex` can take an additional parameter, `fill_value`.

```
In [55]: obj.reindex(['a', 'b', 'c', 'd', 'e'], fill_value=0)
```

```
Out[55]: a   -5.3
         b    7.2
         c    3.6
         d    4.5
         e    0.0
         dtype: float64
```

A different approach uses a `method` parameter that attempts to extrapolate existing data into the new Index. One such method is `ffill`, which “step-fills” the existing data forward. Alternatively, `bfill` “step-fills” the data backwards.

```
In [56]: obj3 = Series(['blue', 'purple', 'yellow'], index=[0, 2, 4])
         obj3.reindex(range(6), method='ffill')
```

```
Out[56]: 0    blue
         1    blue
         2   purple
         3   purple
         4   yellow
         5   yellow
         dtype: object
```

The `reindex` method works for DataFrame objects as well. For DataFrame objects, `reindex` can also specify column reindexing.

```
In [57]: frame = DataFrame(np.arange(9).reshape((3, 3)), index=['a', 'c', 'd'],
                           columns=['Ohio', 'Texas', 'California'])
         frame
```

```
Out[57]:   Ohio  Texas  California
a      0      1           2
c      3      4           5
d      6      7           8
```

```
In [58]: frame.reindex(index=['a', 'b', 'c', 'd'], method='ffill',
                        columns=states)
```

```
Out[58]:
```

	California	Ohio	Oregon	Texas
a	2	0	NaN	1
b	2	0	NaN	1
c	5	3	NaN	4
d	8	6	NaN	7

Alternatively, you can use `ix` to achieve the same effect more concisely.

```
In [59]: frame.ix[['a', 'b', 'c', 'd'], states]
```

```
Out[59]:
```

	California	Ohio	Oregon	Texas
a	2	0	NaN	1
b	NaN	NaN	NaN	NaN
c	5	3	NaN	4
d	8	6	NaN	7

3.2 Dropping entries from an axis

Suppose you have a `Series` object with data you wish to remove. Using the `drop` method, you can specify an index element to remove.

```
In [60]: obj = Series(np.arange(5.), index=['a', 'b', 'c', 'd', 'e'])
         new_obj = obj.drop('c')
         new_obj
```

```
Out[60]: a    0
         b    1
         d    3
         e    4
         dtype: float64
```

You can also drop a list of index elements at once.

```
In [61]: obj.drop(['d', 'c'])
```

```
Out[61]: a    0
         b    1
         e    4
         dtype: float64
```

The same works for `DataFrame` objects and the `drop` method.

```
In [62]: data = DataFrame(np.arange(16).reshape((4, 4)),
                          index=['Ohio', 'Colorado', 'Utah', 'New York'],
                          columns=['one', 'two', 'three', 'four'])
```

```
In [63]: data.drop(['Colorado', 'Ohio'])
```

```
Out[63]:
```

	one	two	three	four
Utah	8	9	10	11
New York	12	13	14	15

Additionally, `DataFrame.drop()` can remove columns by specifying an `axis` parameter.

```
In [64]: data.drop('two', axis=1)
```

```
Out[64]:
```

	one	three	four
Ohio	0	2	3
Colorado	4	6	7
Utah	8	10	11
New York	12	14	15

```
In [65]: data.drop(['two', 'four'], axis=1)
```

```
Out[65]:
```

	one	three
Ohio	0	2
Colorado	4	6
Utah	8	10
New York	12	14

3.3 Indexing, selection, and filtering

In this section we will explore the various techniques available for slicing `Series` and `DataFrame` objects. On the one hand, we can deal with these objects as `dict` structures, accessing elements by requesting their index keys. On the other hand, we can treat these objects as list structures, accessing elements by the order of the index list.

```
In [66]: obj = Series(np.arange(4.), index=['a', 'b', 'c', 'd'])
         obj['b']
```

```
Out[66]: 1.0
```

```
In [67]: obj[1]
```

```
Out[67]: 1.0
```

This flexibility allows you to incorporate all of the previous array slicing that worked for NumPy `ndarray` objects.

```
In [68]: obj[2:4]
```

```
Out[68]: c    2
         d    3
         dtype: float64
```

Conversely, you can use a list of `dict` keys to achieve the same end.

```
In [69]: obj[['b', 'a', 'd']]
```

```
Out[69]: b    1
         a    0
         d    3
         dtype: float64
```

Here are some alternative slicing techniques for `Series` objects.

```
In [70]: obj[[1, 3]]
```

```
Out[70]: b    1
         d    3
         dtype: float64
```

```
In [71]: obj[obj < 2]
```

```
Out[71]: a    0
         b    1
         dtype: float64
```

```
In [72]: obj['b':'c']
```

```
Out[72]: b    1
         c    2
         dtype: float64
```

You can assign values to sub-objects which then reflect on the original object.

```
In [73]: obj['b':'c'] = 5
         obj
```

```
Out[73]: a    0
         b    5
         c    5
         d    3
         dtype: float64
```

The same capabilities are extended to the `DataFrame` objects. The added flexibility is that the same indexing techniques also apply to column slicing as well as index slicing.

```
In [74]: data = DataFrame(np.arange(16).reshape((4, 4)),
                          index=['Ohio', 'Colorado', 'Utah', 'New York'],
                          columns=['one', 'two', 'three', 'four'])

data
```

```
Out[74]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
In [75]: data['two']
```

```
Out[75]: Ohio    1
         Colorado  5
         Utah     9
         New York 13
         Name: two, dtype: int64
```

```
In [76]: data[['three', 'one']]
```

```
Out[76]:
```

	three	one
Ohio	2	0
Colorado	6	4
Utah	10	8
New York	14	12

The natural slicing will always refer to the index list, not the column list, which is useful to keep in mind.

```
data[:2]
```

```
In [77]: data[data['three'] > 5]
```

```
Out[77]:
```

	one	two	three	four
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

Recall that you can generate a corresponding Boolean array by subjecting a `DataFrame` to a boolean statement, such as the following:

```
In [78]: data < 5
```

```
Out[78]:
```

	one	two	three	four
Ohio	True	True	True	True
Colorado	True	False	False	False
Utah	False	False	False	False
New York	False	False	False	False

You can use Boolean arrays to do simple thresholding to your data. You can isolate entries in your data subject to identical Boolean conditions, and manipulate these specific subsets of the data.

```
In [79]: data[data < 5] = 0
data
```

```
Out[79]:
```

	one	two	three	four
Ohio	0	0	0	0
Colorado	0	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

The `DataFrame.ix` field gives you even more powerful ways to slice your data. In general, slicing works by providing two arguments, an index and a column specification, and it will then return that particular subset.

```
In [80]: data.ix['Colorado', ['two', 'three']]
```

```
Out[80]: two      5
         three    6
         Name: Colorado, dtype: int64
```

You can overload requests by using a list of index or column elements. Additionally, you may reorder the indices or columns in your subset by permuting the order of the specified elements, so long as they exist in the original `DataFrame`.

```
In [81]: data.ix[['Colorado', 'Utah'], [3, 0, 1]]
```

```
Out[81]:
```

	four	one	two
Colorado	7	0	5
Utah	11	8	9

The `ix` approach is very powerful. See if you can work through the mechanics of the next few examples to see just how versatile slicing with `ix` actually is.

```
In [82]: data.ix[2]
```

```
Out[82]: one      8
         two      9
         three    10
         four     11
         Name: Utah, dtype: int64
```



```
In [83]: data.ix[:, 'Utah', 'two']
```

```
Out[83]: Ohio      0
         Colorado   5
         Utah       9
         Name: two, dtype: int64
```

```
In [84]: data.ix[data.three > 5, :3]
```

```
Out[84]:
```

	one	two	three
Colorado	0	5	6
Utah	8	9	10
New York	12	13	14

3.4 Arithmetic and data alignment

As we mentioned before, we can do arithmetic on `Series` and `DataFrame` objects.

```
In [85]: s1 = Series([7.3, -2.5, 3.4, 1.5], index=['a', 'c', 'd', 'e'])
         s2 = Series([-2.1, 3.6, -1.5, 4, 3.1], index=['a', 'c', 'e', 'f', 'g'])
         print s1, "\n\n"
         print s2
```

```
a    7.3
c   -2.5
d    3.4
e    1.5
dtype: float64
```

```
a   -2.1
c    3.6
e   -1.5
f    4.0
g    3.1
dtype: float64
```

Importantly, arithmetic is only performed on elements sharing an index. If either object has an index value that the other does not, the arithmetic operation is undefined, so the resultant object contains an `NaN` element.

```
In [86]: s1 + s2
```

```
Out[86]: a    5.2
         c    1.1
         d   NaN
         e    0.0
         f   NaN
         g   NaN
         dtype: float64
```

The same holds for `DataFrame` arithmetic, except now it requires that both the index and column of each `DataFrame` object is well-defined.

```
In [87]: df1 = DataFrame(np.arange(9.).reshape((3, 3)), columns=list('bcd'),
                        index=['Ohio', 'Texas', 'Colorado'])
```

```
df2 = DataFrame(np.arange(12.).reshape((4, 3)), columns=list('bde'),
                index=['Utah', 'Ohio', 'Texas', 'Oregon'])
print df1, "\n\n"
print df2
```

```
b  c  d
Ohio    0  1  2
Texas   3  4  5
Colorado 6  7  8
```

```
      b  d  e
Utah   0  1  2
Ohio   3  4  5
Texas  6  7  8
Oregon 9 10 11
```

```
In [88]: df1 + df2
```

```
Out[88]:
```

	b	c	d	e
Colorado	NaN	NaN	NaN	NaN
Ohio	3	NaN	6	NaN
Oregon	NaN	NaN	NaN	NaN
Texas	9	NaN	12	NaN
Utah	NaN	NaN	NaN	NaN

3.4.1 Arithmetic methods with fill values

Often NaN values are undesirable, as they can cause errors when doing arithmetic operations on the data.

```
In [89]: df1 = DataFrame(np.arange(12.).reshape((3, 4)), columns=list('abcd'))
df2 = DataFrame(np.arange(20.).reshape((4, 5)), columns=list('abcde'))

print df1, "\n\n"
print df2
```

```
a  b  c  d
0  0  1  2  3
1  4  5  6  7
2  8  9 10 11
```

```
      a  b  c  d  e
0  0  1  2  3  4
1  5  6  7  8  9
2 10 11 12 13 14
3 15 16 17 18 19
```

```
In [90]: df1 + df2
```

```
Out[90]:
```

	a	b	c	d	e
0	0	2	4	6	NaN
1	9	11	13	15	NaN
2	18	20	22	24	NaN
3	NaN	NaN	NaN	NaN	NaN

This can be avoided by using the built-in `DataFrame.add()` method, which takes as parameters a `DataFrame` object *and* an optional `fill_value` which deals with otherwise NaN entries.

```
In [91]: df1.add(df2, fill_value=0)
```

```
Out[91]:
```

	a	b	c	d	e
0	0	2	4	6	4
1	9	11	13	15	9
2	18	20	22	24	14
3	15	16	17	18	19

In fact, most `DataFrame` organization methods take `fill_value` as a parameter to deal with undefined cases, such as `reindex`.

```
In [92]: df1.reindex(columns=df2.columns, fill_value=0)
```

```
Out[92]:
```

	a	b	c	d	e
0	0	1	2	3	0
1	4	5	6	7	0
2	8	9	10	11	0

3.5 Operations between DataFrame and Series

Broadcasting NumPy arrays is a very useful technique for performing arithmetic operations concisely. And efficiently, actually. This is because while normal Python arithmetic is *interpreted*, NumPy arithmetic is based on *compiled C* code, which is much more efficient in general.

```
In [93]: arr = np.arange(12.).reshape((3, 4))
         arr
```

```
Out[93]: array([[ 0.,  1.,  2.,  3.],
                [ 4.,  5.,  6.,  7.],
                [ 8.,  9., 10., 11.]])
```

Normally we think of broadcasting a scalar element onto a one-dimensional array vector. In fact, broadcasting is much more powerful, because you can broadcast an *array* over a bigger array.

```
In [94]: arr[0]
```

```
Out[94]: array([ 0.,  1.,  2.,  3.] )
```

```
In [95]: arr - arr[0]
```

```
Out[95]: array([[ 0.,  0.,  0.,  0.],
                [ 4.,  4.,  4.,  4.],
                [ 8.,  8.,  8.,  8.]])
```

`DataFrame` and `Series` objects work along similar lines.

```
In [96]: frame = DataFrame(np.arange(12.).reshape((4, 3)), columns=list('bde'),
                          index=['Utah', 'Ohio', 'Texas', 'Oregon'])
         series = frame.ix[0]
         print frame, "\n\n"
         print series
```

b	d	e
Utah	0	1 2
Ohio	3	4 5
Texas	6	7 8
Oregon	9	10 11

```

b    0
d    1
e    2
Name: Utah, dtype: float64

```

You can broadcast the values in a `Series` over its parent `DataFrame` as you would with NumPy `ndarrays`.

```
In [97]: frame - series
```

```

Out[97]:
      b  d  e
Utah   0  0  0
Ohio   3  3  3
Texas  6  6  6
Oregon  9  9  9

```

Of course, if either a `Series` and `DataFrame` object has index or column values the other does not, the undefined arithmetic simply is sent to `NaN`. (We discussed ways to avoid this issue in the previous sections).

```
In [98]: series2 = Series(range(3), index=['b', 'e', 'f'])
        frame + series2
```

```

Out[98]:
      b  d  e  f
Utah   0 NaN  3 NaN
Ohio   3 NaN  6 NaN
Texas  6 NaN  9 NaN
Oregon  9 NaN 12 NaN

```

```
In [99]: series3 = frame['d']
```

```

print frame, "\n\n"
print series3

```

b	d	e
Utah	0	1 2
Ohio	3	4 5
Texas	6	7 8
Oregon	9	10 11

```

Utah      1
Ohio      4
Texas     7
Oregon    10
Name: d, dtype: float64

```

Using the built-in `DataFrame` arithmetic operations such as `add` or `sub` gives the option to specify the axis (0: index, 1: columns) over which the arithmetic will take place (again, you can use `fill_value` to avoid potential `NaN` values).

```
In [100]: frame.sub(series3, axis=0)
```

```
Out[100]:
```

	b	d	e
Utah	-1	0	1
Ohio	-1	0	1
Texas	-1	0	1
Oregon	-1	0	1

3.6 Function application and mapping

One of the most important capabilities of `Series` and `DataFrame` is the ability to apply function transformations to the data. Every `ufunc` defined by NumPy can be applied to a `DataFrame` (or `Series`) object.

```
In [101]: frame = DataFrame(np.random.randn(4, 3), columns=list('bde'),
                             index=['Utah', 'Ohio', 'Texas', 'Oregon'])

frame
```

```
Out[101]:
```

	b	d	e
Utah	0.450554	0.092673	1.248133
Ohio	0.768101	1.248804	0.774191
Texas	-0.319657	-0.624964	1.078814
Oregon	0.544647	0.855588	1.343268

For example, you can apply a nonnegativity transform by including a built-in NumPy absolute value.

```
In [102]: np.abs(frame)
```

```
Out[102]:
```

	b	d	e
Utah	0.450554	0.092673	1.248133
Ohio	0.768101	1.248804	0.774191
Texas	0.319657	0.624964	1.078814
Oregon	0.544647	0.855588	1.343268

You can define and apply custom functions in two fashions. One is by using lambdas to construct anonymous functions:

```
In [103]: frame.apply(lambda x: x.max() - x.min())
```

```
Out[103]: b    1.087758
          d    1.873768
          e    0.569077
          dtype: float64
```

```
In [104]: frame.apply(lambda x: x.max() - x.min(), axis=1)
```

```
Out[104]: Utah    1.155460
          Ohio    0.480703
          Texas    1.703778
          Oregon    0.798621
          dtype: float64
```

Alternatively, you can define your own unary function and simply apply it using the same overall approach.

```
In [105]: def f(x):
           return Series([x.min(), x.max()], index=['min', 'max'])
frame.apply(f)
```

```
Out[105]:
```

	b	d	e
min	-0.319657	-0.624964	0.774191
max	0.768101	1.248804	1.343268

For presentations and general readability, it is useful to format decimal or date values into condensed forms, and **Pandas** lets you achieve this by using the `applymap` method for **DataFrame** and **Series** objects. The difference between `apply` and `applymap` is rather subtle and often functionally negligible, but the idea is that `apply` works on a particular subsets of rows or columns, whereas `applymap` is element-wise.

```
In [106]: format = lambda x: '%.2f' % x
          frame.applymap(format)
```

```
Out[106]:
```

	b	d	e
Utah	0.45	0.09	1.25
Ohio	0.77	1.25	0.77
Texas	-0.32	-0.62	1.08
Oregon	0.54	0.86	1.34

Alternatively, you can use the built-in Python `map` function.

```
In [107]: frame['e'].map(format)
```

```
Out[107]:
```

Utah	1.25
Ohio	0.77
Texas	1.08
Oregon	1.34

Name: e, dtype: object

Exercise: Determine which of `apply` or `map` is computationally more efficient.

3.7 Sorting and ranking

One fundamental problem in data analysis, let alone computer science in general, is sorting data. **Pandas** provides a number of techniques for sorting information in the index, columns, and the actual data itself.

The first technique is `sort_index`, which is a method for both **Series** and **DataFrame** objects. For **Series** objects, `sort_index` works as follows:

```
In [108]: obj = Series(range(4), index=['d', 'a', 'b', 'c'])
          obj.sort_index()
```

```
Out[108]:
```

a	1
b	2
c	3
d	0

dtype: int64

Since there is only one meaningful index to sort, the labels, `sort_index` is a very intuitive method. I want to point out that `sort_index` does not have side-effects; that is, calling `sort_index` on an object does not actually change the internals of the object itself. Instead, a sorted copy of the original object is produced.

The method `sort_index` works similarly with **DataFrame** objects, but now there are two potential axes along which to sort. The default is the `index`, as we see below:

```
In [109]: frame = DataFrame(np.arange(8).reshape((2, 4)), index=['three', 'one'],
                           columns=['d', 'a', 'b', 'c'])
          frame.sort_index()
```

```
Out[109]:
```

	d	a	b	c
one	4	5	6	7
three	0	1	2	3

By specifying the axis as a parameter, one can choose the columns instead. (Recall that in Python everything begins at 0, so the second axis corresponds to axis number 1).

```
In [110]: frame.sort_index(axis=1)
```

```
Out[110]:
```

	a	b	c	d
three	1	2	3	0
one	5	6	7	4

The `sort_index` method also allows you to flip the ordering by specifying the `ascending` parameter.

```
In [111]: frame.sort_index(axis=1, ascending=False)
```

```
Out[111]:
```

	d	c	b	a
three	0	3	2	1
one	4	7	6	5

If you want to sort the elements themselves, as opposed to the index, Pandas provides the `order` method for `Series` objects.

```
In [112]: obj = Series([4, 7, -3, 2])
          obj.order()
```

```
Out[112]:
```

2	-3
3	2
0	4
1	7

dtype: int64

By default, NaN values are placed at the end upon sorting the `Series`.

```
In [113]: obj = Series([4, np.nan, 7, np.nan, -3, 2])
          obj.order()
```

```
Out[113]:
```

4	-3
5	2
0	4
2	7
1	NaN
3	NaN

dtype: float64

For `DataFrame` objects you can specify the index or column you wish to sort. Additionally, if your data set is properly constructed, you can sort by two columns or indices, as the below example exhibits:

```
In [114]: frame = DataFrame({'b': [4, 7, -3, 2], 'a': [0, 1, 0, 1]})
```

```
frame.sort_index(by='b')
```

```
Out[114]:
```

	a	b
2	0	-3
3	1	2
0	0	4
1	1	7

```
In [115]: frame.sort_index(by=['a', 'b'])
```

```
Out[115]:
```

	a	b
2	0	-3
0	0	4
3	1	2
1	1	7

3.8 Axis indexes with duplicate values

It is possible for a `Series` or `DataFrame` object not to have a unique index. For example:

```
In [116]: obj = Series(range(5), index=['a', 'a', 'b', 'b', 'c'])
          obj
```

```
Out[116]: a    0
          a    1
          b    2
          b    3
          c    4
          dtype: int64
```

Pandas has a field for the index of any object to indicate whether or not the index is unique (no duplicate indices).

```
In [117]: obj.index.is_unique
```

```
Out[117]: False
```

If an index is not unique, then slicing the object for a repeated index returns a sub-object. For example:

```
In [118]: obj['a']
```

```
Out[118]: a    0
          a    1
          dtype: int64
```

For unique index items, the default return-type is a scalar:

```
In [119]: obj['c']
```

```
Out[119]: 4
```

The same goes for `DataFrame` objects, although they are more complicated.

```
In [120]: df = DataFrame(np.random.randn(4, 3), index=['a', 'a', 'b', 'b'])
          df
```

```
Out[120]:
```

	0	1	2
a	-0.267175	1.793095	-0.652929
a	-1.886837	1.059626	0.644448
b	-0.007799	-0.449204	2.448963
b	0.667226	0.802926	0.575721

Remember that you have to slice the index using `ix`, and you will observe the same behavior.

```
In [121]: df.ix['b']
```

```
Out[121]:
```

	0	1	2
b	-0.007799	-0.449204	2.448963
b	0.667226	0.802926	0.575721

4 Summarizing and computing descriptive statistics

Oftentimes, you need a quick way to come up with basic summary statistics of data sets. The solution that Pandas provides is incredibly robust, especially with regard to NaN entries.

```
In [122]: df = DataFrame([[1.4, np.nan], [7.1, -4.5],
                        [np.nan, np.nan], [0.75, -1.3]],
                        index=['a', 'b', 'c', 'd'],
                        columns=['one', 'two'])

df
```

```
Out[122]:
```

	one	two
a	1.40	NaN
b	7.10	-4.5
c	NaN	NaN
d	0.75	-1.3

By default, the `sum` method will skip NaN entries for each column in a `DataFrame`.

```
In [123]: df.sum()
```

```
Out[123]: one      9.25
         two     -5.80
         dtype: float64
```

For the `DataFrame` object, you can also apply along either the index axis or the column axis. Again, `sum` will skip over NaN elements when arriving at a value.

```
In [124]: df.sum(axis=1)
```

```
Out[124]: a      1.40
         b      2.60
         c      NaN
         d     -0.55
         dtype: float64
```

If you don't want this behavior, you can always tell the statistics function you are applying not to skip the NaN entries. Here is an example using `mean`:

```
In [125]: df.mean(axis=1, skipna=False)
```

```
Out[125]: a      NaN
         b      1.300
         c      NaN
         d     -0.275
         dtype: float64
```

Another useful statistic is `idxmax`, which returns the index of the maximum value of a column in a `DataFrame`.

```
In [126]: df.idxmax()
```

```
Out[126]: one      b
         two      d
         dtype: object
```

One incredibly useful method is `cumsum`, which has a number of important applications in the analysis of probability distributions and random walks.

```
In [127]: df.cumsum()
```

```
Out[127]:
```

	one	two
a	1.40	NaN
b	8.50	-4.5
c	NaN	NaN
d	9.25	-5.8

You can also get a quick overview of all of the summary statistics of a **DataFrame** simply by calling the `describe` method.

```
In [128]: df.describe()
```

```
Out[128]:
```

	one	two
count	3.000000	2.000000
mean	3.083333	-2.900000
std	3.493685	2.262742
min	0.750000	-4.500000
25%	1.075000	-3.700000
50%	1.400000	-2.900000
75%	4.250000	-2.100000
max	7.100000	-1.300000

Method	Description
count	Number of non-NaN values
describe	Compute set of summary statistics for Series or each DataFrame column
min, max	Compute minimum and maximum values
argmin, argmax	Compute index locations for minimum and maximum values
idmin, idmax	Compute index values for minimum and maximum values
quantile	Compute sample quantile ranging from 0 to 1
sum	Sum of values
mean	Mean of values
median	Arithmetic median of values
mad	Mean absolute deviation from mean value
var	Sample variance of values
std	Sample standard deviation of values
skew	Sample skewness (3rd moment) of values
kurt	Sample kurtosis (4th moment) of values
cumsum	Cumulative sum of values
cummin, cummax	Cumulative min and max of values
cumprod	cumulative product of values
diff	Compute 1st arithmetic difference (useful for time series
pct_change	Compute percent changes

Descriptive and Summary Statistics Series objects also have a `describe` method. The `describe` method outputs statistics based on the `dtype` of the underlying object. In the above example, `df` had a `dtype` of `float64`, so `describe` produced information pertinent to floating-point numerics. In the below example, the Series object has a `dtype` of `object`, which results in different summary statistics.

```
In [129]: obj = Series(['a', 'a', 'b', 'c'] * 4)
          obj.describe()
```

```
Out[129]: count      16
          unique      3
          top         a
          freq        8
          dtype: object
```

4.1 Correlation and covariance

One common problem in data analysis, especially in the analysis of time series data like historical prices for financial securities, is correlation and covariance analysis. To this end **Pandas** has a number of features to make the analysis simple.

Here is one example, using a built-in data aggregator using [Yahoo! Finance](#) in the **Pandas** API. Returns on a stock are defined as the percent change in the stock's closing value from day-to-day.

```
In [130]: import pandas.io.data as web

          all_data = {}
          for ticker in ['AAPL', 'IBM', 'MSFT', 'CSCO']:
              all_data[ticker] = web.get_data_yahoo(ticker)

          price = DataFrame({tic: data['Adj Close']
                             for tic, data in all_data.iteritems()})
          volume = DataFrame({tic: data['Volume']
                              for tic, data in all_data.iteritems()})

          returns = price.pct_change()
          returns.tail()
```

```
Out[130]:
```

	AAPL	CSCO	IBM	MSFT
Date				
2015-03-25	-0.026127	-0.019431	-0.023313	-0.033566
2015-03-26	0.006970	-0.013578	0.008731	-0.006030
2015-03-27	-0.007968	0.001488	-0.001183	-0.005824
2015-03-30	0.025314	0.019316	0.014152	-0.000244
2015-03-31	-0.015352	0.003280	-0.013340	-0.007324

When given a Series object, the `corr` method computes the scalar correlation between the Series and another Series.

```
In [131]: returns.MSFT.corr(returns.IBM)
```

```
Out[131]: 0.50144660652462925
```

By contrast, `corr` and `cov` returns a correlation and covariance matrix **DataFrame** with filled correlation and covariance values, respectively.

```
In [132]: returns.MSFT.cov(returns.IBM)
```

```
Out[132]: 8.3470993744258389e-05
```

```
In [133]: returns.corr()
```

```
Out[133]:
```

	AAPL	CSCO	IBM	MSFT
AAPL	1.000000	0.335746	0.371383	0.347494
CSCO	0.335746	1.000000	0.448426	0.464868
IBM	0.371383	0.448426	1.000000	0.501447
MSFT	0.347494	0.464868	0.501447	1.000000

```
In [134]: returns.cov()
```

```
Out[134]:
```

	AAPL	CSCO	IBM	MSFT
AAPL	0.000283	0.000098	0.000073	0.000083
CSCO	0.000098	0.000301	0.000091	0.000114
IBM	0.000073	0.000091	0.000138	0.000083
MSFT	0.000083	0.000114	0.000083	0.000200

The `corrwith` method computes pairwise correlations and stores the resultant in a `Series`. Note that the correlation between IBM and IBM is 1.

```
In [135]: returns.corrwith(returns.IBM)
```

```
Out[135]: AAPL    0.371383
          CSCO    0.448426
          IBM     1.000000
          MSFT    0.501447
          dtype: float64
```

Passing a `DataFrame` instead computes correlation with like-columns.

```
In [136]: returns.corrwith(volume)
```

```
Out[136]: AAPL    -0.097904
          CSCO    -0.235580
          IBM     -0.184944
          MSFT    -0.120080
          dtype: float64
```

4.2 Unique values, value counts, and membership

Given data with repeats, you can eliminate the excess by using the `unique` method.

```
In [137]: obj = Series(['c', 'a', 'd', 'a', 'a', 'b', 'b', 'c', 'c'])
          uniques = obj.unique()
          uniques
```

```
Out[137]: array(['c', 'a', 'd', 'b'], dtype=object)
```

The `value_counts` returns a `Series` with an index made up of the unique entries in the original `Series`, and the new entries give the total appearances of each value.

```
In [138]: obj.value_counts()
```

```
Out[138]: c      3
          a      3
          b      2
          d      1
          dtype: int64
```

You can perform set-membership operations to, for example, construct masks which you can then apply to your original data.

```
In [139]: mask = obj.isin(['b', 'c']) # This forms a Series object of Boolean values
          obj[mask]
```

```
Out[139]: 0    c
          5    b
          6    b
          7    c
          8    c
          dtype: object
```

For DataFrame objects, you can apply the `value_counts` method to each subseries, producing a new DataFrame of frequency statistics.

```
In [140]: data = DataFrame({'Qu1': [1, 3, 4, 3, 4],
                           'Qu2': [2, 3, 1, 2, 3],
                           'Qu3': [1, 5, 2, 4, 4]})
          data.apply(pd.value_counts).fillna(0)
```

```
Out[140]:   Qu1  Qu2  Qu3
          1    1    1    1
          2    0    2    1
          3    2    2    0
          4    2    0    2
          5    0    0    1
```

5 Handling missing data

One of the primary problems with data analysis is the prevalence of missing data. In many cases, arithmetic operations, summary statistics, and other functions require that your data be intact in order to provide meaningful results. Pandas gives a number of functions to address the problem of missing data, allowing you to filter it out easily.

Consider this Series of string values.

```
In [141]: string_data = Series(['aardvark', 'artichoke', np.nan, 'avocado'])
          string_data
```

```
Out[141]: 0    aardvark
          1    artichoke
          2         NaN
          3     avocado
          dtype: object
```

The `isnull` method identifies every NaN entry. Alternatively, `notnull` will identify every non-NaN entry.

```
In [142]: string_data.isnull()
```

```
Out[142]: 0    False
          1    False
          2     True
          3    False
          dtype: bool
```

```
In [143]: string_data[0] = None
          string_data.notnull()
```

```
Out[143]: 0    False
          1     True
          2    False
          3     True
          dtype: bool
```

5.1 Filtering out missing data

A simple way to remove missing entries from a `Series` object is to use `dropna`.

```
In [144]: from numpy import nan as NA
          data = Series([1, NA, 3.5, NA, 7])
          data.dropna()
```

```
Out[144]: 0    1.0
          2    3.5
          4    7.0
          dtype: float64
```

Alternatively, you can use `Boolean Series` and `notnull` to mask the original data.

```
In [145]: data[data.notnull()]
```

```
Out[145]: 0    1.0
          2    3.5
          4    7.0
          dtype: float64
```

`DataFrame` objects are trickier. For example, how should `Pandas` handle a mostly-complete row? The correct answer is ambiguous. By default, `dropna` will eliminate *any* row with a `NaN` (we redefined `NaN` to `NA` here) value.

```
In [146]: data = DataFrame([[1., 6.5, 3.], [1., NA, NA],
                           [NA, NA, NA], [NA, 6.5, 3.]])
          cleaned = data.dropna()
          data
```

```
Out[146]:    0    1    2
0    1  6.5    3
1    1  NaN  NaN
2  NaN  NaN  NaN
3  NaN  6.5    3
```

```
In [147]: cleaned
```

```
Out[147]:    0    1    2
0    1  6.5    3
```

Alternatively, you can require that a row be eliminated only if it is *completely* empty.

```
In [148]: data.dropna(how='all')
```

```
Out[148]:    0    1    2
0    1  6.5    3
1    1  NaN  NaN
3  NaN  6.5    3
```

You can also specify columns for deletion. Again, you can change the deletion requirements as needed.

```
In [149]: data[4] = NA # fill a column entirely with NA
          data.dropna(axis=1, how='all')
```

```
Out[149]:
```

	0	1	2
0	0	1	6.5
1	1	1	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3

The `dropna` method is very robust. You can also specify a minimum threshold of data in a particular row as a criterion for deletion. In the next example, we threshold at 2 entries per row, allowing rows with one NaN value to stay while deleting any more patchy rows.

```
In [150]: df = DataFrame(np.random.randn(7, 3))
          df.ix[:4, 1] = NA; df.ix[:2, 2] = NA
          df
```

```
Out[150]:
```

	0	1	2
0	1.381918	NaN	NaN
1	-0.206282	NaN	NaN
2	1.811659	NaN	NaN
3	-1.313554	NaN	-0.615939
4	0.174072	NaN	-0.035408
5	-1.194063	1.037339	-0.364987
6	-0.215726	0.763525	0.671308

```
In [151]: df.dropna(thresh=2)
```

```
Out[151]:
```

	0	1	2
3	-1.313554	NaN	-0.615939
4	0.174072	NaN	-0.035408
5	-1.194063	1.037339	-0.364987
6	-0.215726	0.763525	0.671308

5.2 Filling in missing data

Instead of eliminating missing data outright, Pandas lets you fill in the missing values. The simple approach, using `fillna`, is to pass a value that will then replace every NaN entry.

```
In [152]: df.fillna(0)
```

```
Out[152]:
```

	0	1	2
0	1.381918	0.000000	0.000000
1	-0.206282	0.000000	0.000000
2	1.811659	0.000000	0.000000
3	-1.313554	0.000000	-0.615939
4	0.174072	0.000000	-0.035408
5	-1.194063	1.037339	-0.364987
6	-0.215726	0.763525	0.671308

Alternatively, you can specify different fill values in different columns by giving a `dict` with keys of column names.

```
In [153]: df.fillna({1: 0.5, 2: -1})
```

```
Out[153]:
```

	0	1	2
0	1.381918	0.500000	-1.000000
1	-0.206282	0.500000	-1.000000
2	1.811659	0.500000	-1.000000
3	-1.313554	0.500000	-0.615939
4	0.174072	0.500000	-0.035408
5	-1.194063	1.037339	-0.364987
6	-0.215726	0.763525	0.671308

Using the `inplace` argument, you can overwrite the original `DataFrame` object.

```
In [154]: # always returns a reference to the filled object
_ = df.fillna(0, inplace=True)
df
```

```
Out[154]:
```

	0	1	2
0	1.381918	0.000000	0.000000
1	-0.206282	0.000000	0.000000
2	1.811659	0.000000	0.000000
3	-1.313554	0.000000	-0.615939
4	0.174072	0.000000	-0.035408
5	-1.194063	1.037339	-0.364987
6	-0.215726	0.763525	0.671308

The other main filling technique is to fill by procedure. `ffill` will copy the previous value in a column into the NaN entry.

```
In [155]: df = DataFrame(np.random.randn(6, 3))
df.ix[2::2, 1] = NA; df.ix[4:, 2] = NA
print df, "\n\n"
print df.fillna(method='ffill')
```

```
0      1      2
0 -1.810961 -0.246414 -0.205597
1  0.998181  0.625001  0.410271
2  0.063753      NaN  0.289051
3 -2.202882 -0.068072  0.033761
4  1.840764      NaN      NaN
5 -0.024116  1.462648      NaN
```

```
0      1      2
0 -1.810961 -0.246414 -0.205597
1  0.998181  0.625001  0.410271
2  0.063753  0.625001  0.289051
3 -2.202882 -0.068072  0.033761
4  1.840764 -0.068072  0.033761
5 -0.024116  1.462648  0.033761
```

In cases where you don't want this to extend indefinitely, you can limit the fill method to a certain number of NaN entries after the last available one.

```
In [156]: df.fillna(method='ffill', limit=1)
```

```
Out[156]:
```

	0	1	2
0	-1.810961	-0.246414	-0.205597


```

1  0.998181  0.625001  0.410271
2  0.063753  0.625001  0.289051
3 -2.202882 -0.068072  0.033761
4  1.840764 -0.068072  0.033761
5 -0.024116  1.462648      NaN

```

6 Hierarchical indexing

From “Python for Data Analysis”:

Hierarchical indexing is an important feature of pandas enabling you to have multiple (two or more) index *levels* on an axis. Somewhat abstractly, it provides a way for you to work with higher dimensional data in a lower dimensional form. Let’s start with a simple example; create a **Series** with a list of lists or arrays as the index:

```

In [157]: data = Series(np.random.randn(10),
                        index=[['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'd', 'd'],
                              [1, 2, 3, 1, 2, 3, 1, 2, 2, 3]])

```

data

```

Out[157]: a 1    0.221755
          2   -0.838210
          3    1.396553
         b 1   -1.553775
          2   -0.007680
          3    1.335753
         c 1   -1.296638
          2    1.067990
         d 2   -0.743429
          3    0.500286
dtype: float64

```

Notice how the index **a** corresponds to the sub-indices 1 and 2, and their corresponding data. The **index** object is thus not a simple list but a series of lists corresponding to the inner sub-indices.

```

In [158]: data.index

```

```

Out[158]: MultiIndex(levels=[['a', 'b', 'c', 'd'], [1, 2, 3]],
                      labels=[[0, 0, 0, 1, 1, 1, 2, 2, 3, 3], [0, 1, 2, 0, 1, 2, 0, 1, 1, 2]])

```

Accessing an outer label will give you the sub-**Series** that it corresponds to.

```

In [159]: print data['b'], "\n\n"
          print data['b':'c']

```

```

1   -1.553775
2   -0.007680
3    1.335753
dtype: float64

```

```

b 1   -1.553775
   2   -0.007680
   3    1.335753
c 1   -1.296638
   2    1.067990
dtype: float64

```

You can access sub-indices, which returns the **Series** of all upper indices and their corresponding values.

```
In [160]: data[:, 2]
```

```
Out[160]: a    -0.838210
          b    -0.007680
          c     1.067990
          d    -0.743429
          dtype: float64
```

You can use `unstack` to take the multi-index and place it into a **DataFrame** object.

```
In [161]: data.unstack()
```

```
Out[161]:
```

	1	2	3
a	0.221755	-0.838210	1.396553
b	-1.553775	-0.007680	1.335753
c	-1.296638	1.067990	NaN
d	NaN	-0.743429	0.500286

The inverse of `unstack` is `stack`. Observe:

```
In [162]: data.unstack().stack()
```

```
Out[162]: a  1    0.221755
          2   -0.838210
          3    1.396553
          b  1   -1.553775
          2   -0.007680
          3    1.335753
          c  1   -1.296638
          2    1.067990
          d  2   -0.743429
          3    0.500286
          dtype: float64
```

Multi-indexing has a similar logic with **DataFrame** objects, but it becomes more complicated as both the index and the columns can be given a hierarchy:

```
In [163]: frame = DataFrame(np.arange(12).reshape((4, 3)),
                             index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],
                             columns=[['Ohio', 'Ohio', 'Colorado'],
                                       ['Green', 'Red', 'Green']])

frame
```

```
Out[163]:
```

		Ohio		Colorado
		Green	Red	Green
a	1	0	1	2
	2	3	4	5
b	1	6	7	8
	2	9	10	11

For clarity, let's rename the labels so we know what level we are looking at.

```
In [164]: frame.index.names = ['key1', 'key2']
          frame.columns.names = ['state', 'color']
          frame
```

```
Out[164]: state      Ohio      Colorado
color      Green Red      Green
key1 key2
a      1      0      1      2
      2      3      4      5
b      1      6      7      8
      2      9     10     11
```

Now by specifying any column, whether on the top level or any sublevel, you can get the `DataFrame` of values corresponding to the name.

```
In [165]: frame['Ohio']
```

```
Out[165]: color      Green  Red
key1 key2
a      1      0      1
      2      3      4
b      1      6      7
      2      9     10
```

`MultiIndex` objects are independent in `Pandas`, meaning that you can create them without a corresponding `DataFrame` and reuse them as needed.

```
In [166]: from pandas import MultiIndex
MultiIndex.from_arrays([[ 'Ohio', 'Ohio', 'Colorado'], ['Green', 'Red', 'Green']],
                      names=[ 'state', 'color'])
```

```
Out[166]: MultiIndex(levels=[[u'Colorado', u'Ohio'], [u'Green', u'Red']],
                      labels=[[1, 1, 0], [0, 1, 0]],
                      names=[u'state', u'color'])
```

6.1 Reordering and sorting levels

You can always swap indices on the same level. For example, if you want `key2` and `key1` to switch, you can write

```
In [167]: frame.swaplevel('key1', 'key2')
```

```
Out[167]: state      Ohio      Colorado
color      Green Red      Green
key2 key1
1      a      0      1      2
2      a      3      4      5
1      b      6      7      8
2      b      9     10     11
```

Additionally, you can sort a particular index (in general, you can't sort them all). Specify the index by its order (first is 0, second is 2), and you will see the sort take place:

```
In [168]: frame.sortlevel(1)
```

```
Out[168]: state      Ohio      Colorado
color      Green Red      Green
key1 key2
a      1      0      1      2
b      1      6      7      8
a      2      3      4      5
b      2      9     10     11
```

As with the object-oriented paradigm, you can combine these actions into one statement. For example:

```
In [169]: frame.swaplevel(0, 1).sortlevel(0)
```

```
Out[169]: state      Ohio      Colorado
          color      Green Red      Green
          key2 key1
          1      a          0      1          2
              b          6      7          8
          2      a          3      4          5
              b          9     10         11
```

6.2 Summary statistics by level

With hierarchical indexing, you can specify the level and axis with which to compute summary statistics. If one wants to compute the sum of all values in the `key2` index, you get the relevant sub-DataFrame.

```
In [170]: frame.sum(level='key2')
```

```
Out[170]: state      Ohio      Colorado
          color      Green Red      Green
          key2
          1          6      8          10
          2         12     14          16
```

This of course gets extended to the columns as well, which you have grown accustomed to with `DataFrame` methods.

```
In [171]: frame.sum(level='color', axis=1)
```

```
Out[171]: color      Green      Red
          key1 key2
          a      1          2      1
              2          8      4
          b      1         14      7
              2         20     10
```

6.3 Using a DataFrame's columns

In the examples above we showed how to `stack` and `unstack` `Series` objects into `DataFrames`. But in general `DataFrame` objects give you a lot of discretion regarding which columns you want to convert into indices.

```
In [172]: frame = DataFrame({'a': range(7), 'b': range(7, 0, -1),
                             'c': ['one', 'one', 'one', 'two', 'two', 'two', 'two'],
                             'd': [0, 1, 2, 0, 1, 2, 3]})
```

frame

```
Out[172]:   a  b   c  d
0  0  7  one  0
1  1  6  one  1
2  2  5  one  2
3  3  4  two  0
4  4  3  two  1
5  5  2  two  2
6  6  1  two  3
```

You can overload `set_index` with more than one column to produce a hierarchical index using the values of each respective column.

```
In [173]: frame2 = frame.set_index(['c', 'd'])
          frame2
```

```
Out[173]:
```

	a	b
c d		
one 0	0	7
1 1	1	6
2 2	2	5
two 0	3	4
1 4	4	3
2 5	5	2
3 6	6	1

Crucially, the default **Pandas** behavior is to remove the indexed columns. You can force **Pandas** to keep the old columns by specifying the `drop` parameter:

```
In [174]: frame.set_index(['c', 'd'], drop=False)
```

```
Out[174]:
```

	a	b	c	d
c d				
one 0	0	7	one	0
1 1	1	6	one	1
2 2	2	5	one	2
two 0	3	4	two	0
1 4	4	3	two	1
2 5	5	2	two	2
3 6	6	1	two	3

```
In [175]: frame2.reset_index()
```

```
Out[175]:
```

	c	d	a	b
0	one	0	0	7
1	one	1	1	6
2	one	2	2	5
3	two	0	3	4
4	two	1	4	3
5	two	2	5	2
6	two	3	6	1